



# Database Management Systems

## Module 06: Introduction to SQL/1

**Partha Pratim Das**

*Department of Computer Science and Engineering  
Indian Institute of Technology, Kharagpur*

ppd@cse.iitkgp.ernet.in

**Srijoni Majumdar  
Himadri B G S Bhuyan  
Gurunath Reddy M**



**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
[www.db-book.com](http://www.db-book.com)

Slides marked with 'PPD' are new or edited



# Week 01 Recap

- **Module 01: Course Overview**
  - Why Databases?
  - KYC: Know Your Course
- **Module 02: Introduction to DBMS/1**
  - Levels of Abstraction
  - Schema & Instance
  - Data Models
  - DDL & DML
  - SQL
  - Database Design
- **Module 03: Introduction to DBMS/2**
  - Database Design
  - Database Engine, Users, Architecture
  - History of DBMS
- **Module 04: Introduction to Relational Model/1**
  - Attribute Types
  - Relation Schema and Instance
  - Keys
  - Relational Query Languages
- **Module 05: Introduction to Relational Model/2**
  - Relational Operations
  - Aggregate Operations



# Module Objectives

- To understand relational query language
- To understand data definition and basic query structure



# Module Outline

- History of SQL
- Data Definition Language (DDL)
- Basic Query Structure (DML)



- **History of SQL**
- Data Definition Language (DDL)
- Basic Query Structure (DML)

# HISTORY OF SQL



# History of Query Language

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
  - SQL-86
  - SQL-89
  - SQL-92
  - SQL:1999 (language name became Y2K compliant!)
  - SQL:2003
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
  - Not all examples here may work on your particular system



PPD

- History of SQL
- **Data Definition Language (DDL)**
- Basic Query Structure (DML)

# DATA DEFINITION LANGUAGE (DDL)



# Data Definition Language

The SQL data-definition language (DDL) allows the specification of information about relations, including:

- The schema for each relation
- The domain of values associated with each attribute
- Integrity constraints
- And as we will see later, also other information such as
  - The set of indices to be maintained for each relations
  - Security and authorization information for each relation
  - The physical storage structure of each relation on disk



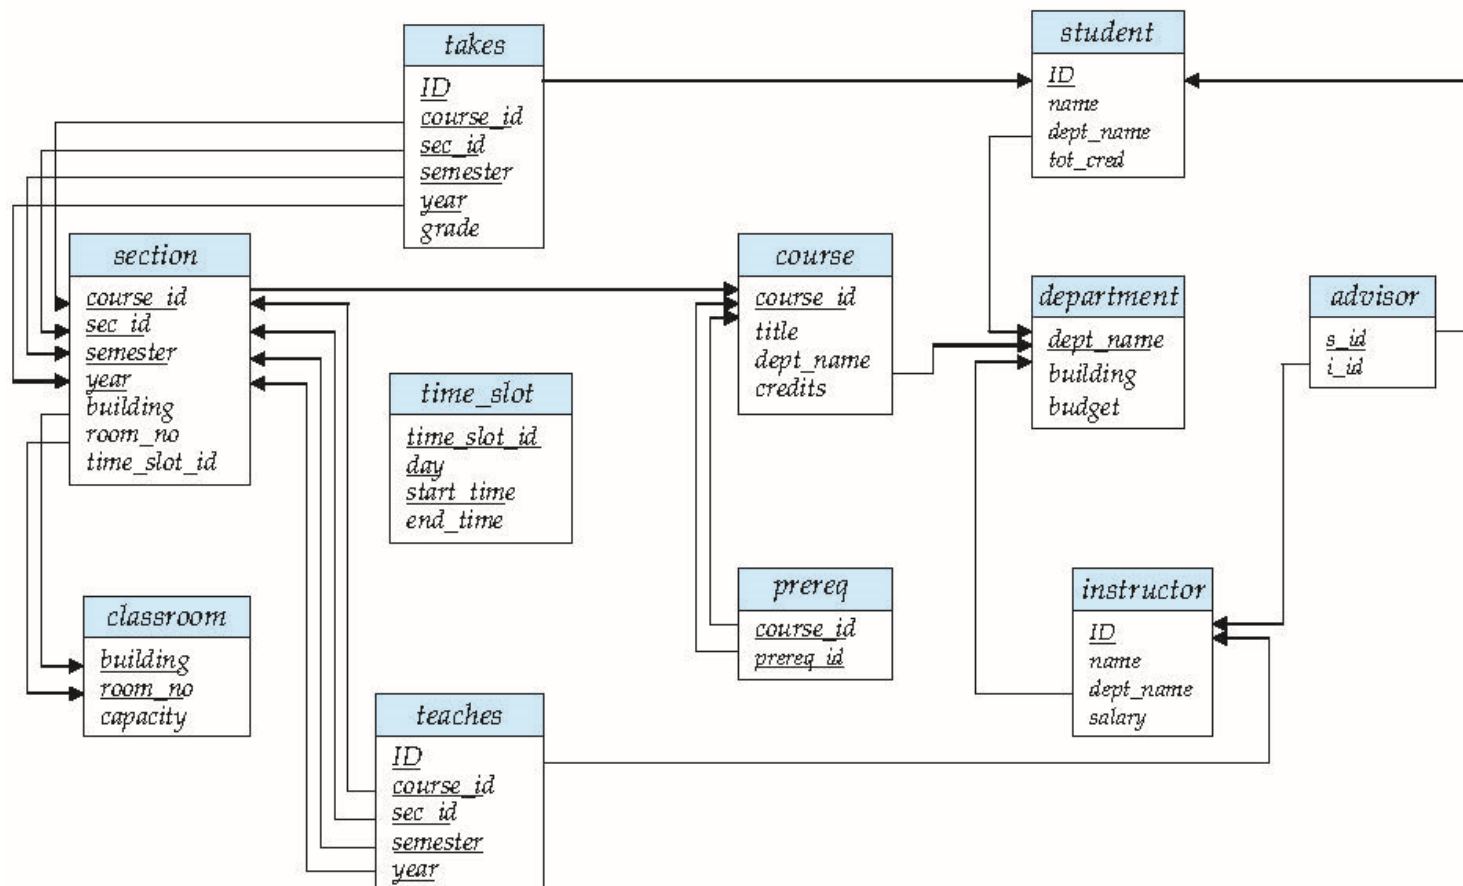


# Domain Types in SQL

- **char(*n*).** Fixed length character string, with user-specified length *n*
- **varchar(*n*).** Variable length character strings, with user-specified maximum length *n*
- **int.** Integer (a finite subset of the integers that is machine-dependent)
- **smallint.** Small integer (a machine-dependent subset of the integer domain type)
- **numeric(*p,d*).** Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point. (ex., **numeric(3,1)**, allows 44.5 to be stores exactly, but not 444.5 or 0.32)
- **real, double precision.** Floating point and double-precision floating point numbers, with machine-dependent precision
- **float(*n*).** Floating point number, with user-specified precision of at least *n* digits
- More are covered in Chapter 4



# Schema Diagram for University Database





# Create Table Construct

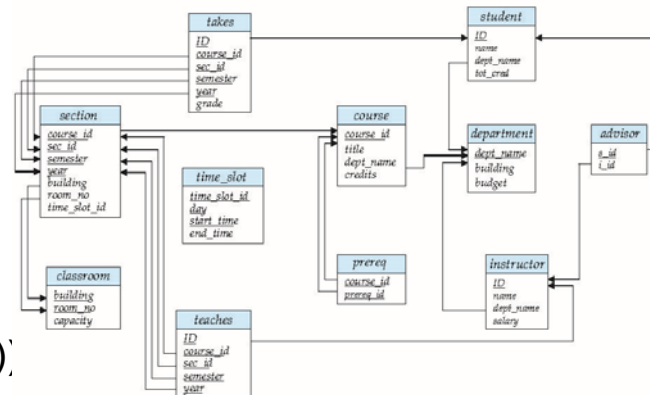
- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1, A2 D2, ..., An Dn,
               (integrity-constraint1),
               ...,
               (integrity-constraintk))
```

- r* is the name of the relation
- each *A<sub>i</sub>* is an attribute name in the schema of relation *r*
- D<sub>i</sub>* is the data type of values in the domain of attribute *A<sub>i</sub>*

- Example:

```
create table instructor (
    ID          char(5),
    name        varchar(20),
    dept_name   varchar(20),
    salary      numeric(8,2)
```





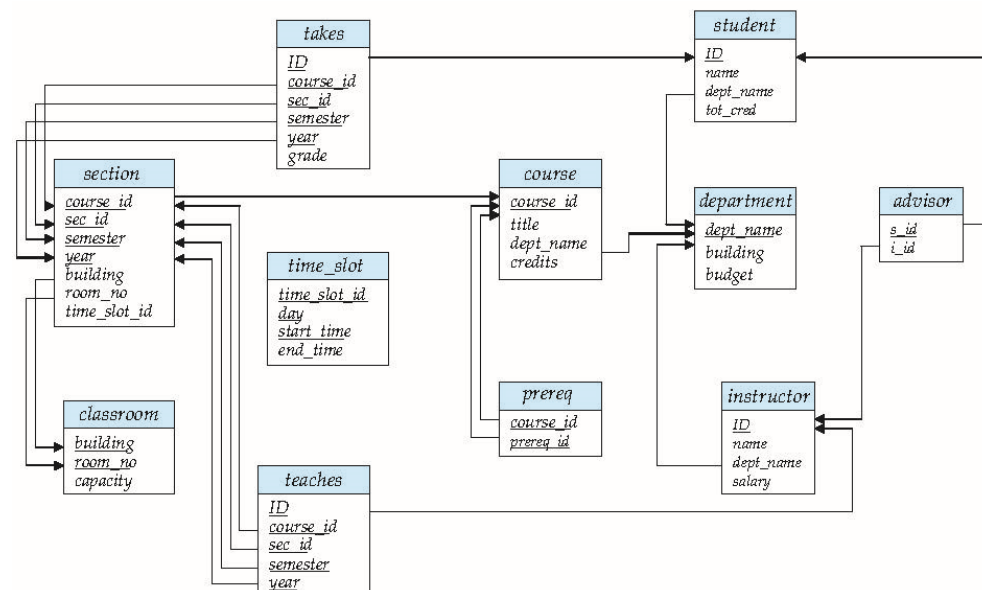
# Integrity Constraints in Create Table

PPD

- **not null**
- **primary key** ( $A_1, \dots, A_n$ )
- **foreign key** ( $A_m, \dots, A_n$ ) **references**  $r$

Example:

```
create table instructor (
    ID          char(5),
    name        varchar(20) not null,
    dept_name   varchar(20),
    salary      numeric(8,2),
    primary key (ID),
    foreign key (dept_name) references department);
```

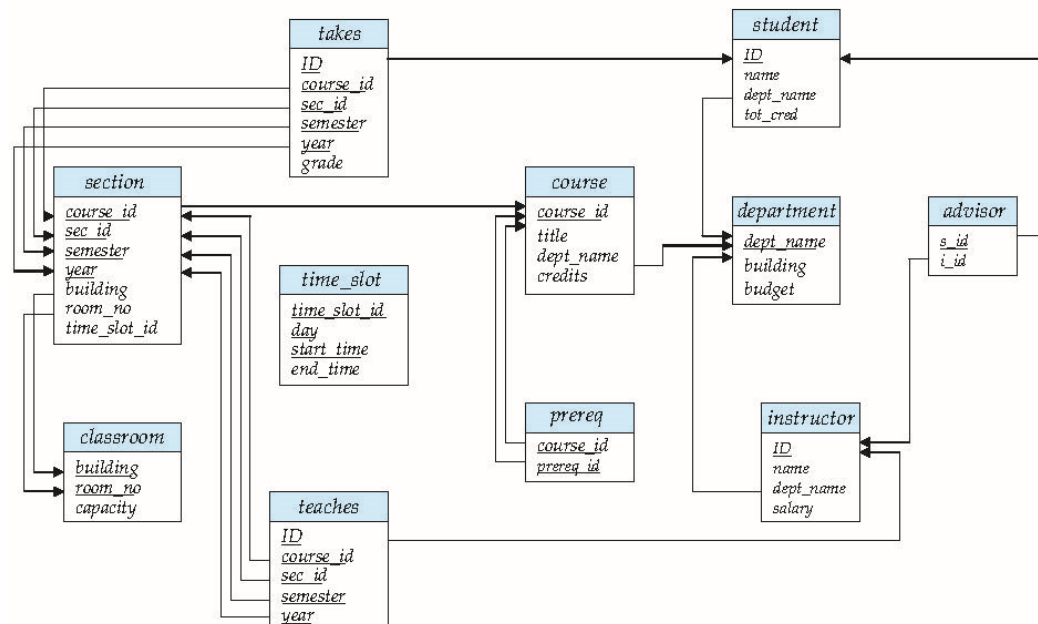


**primary key** declaration on an attribute automatically ensures **not null**



## And a Few More Relation Definitions

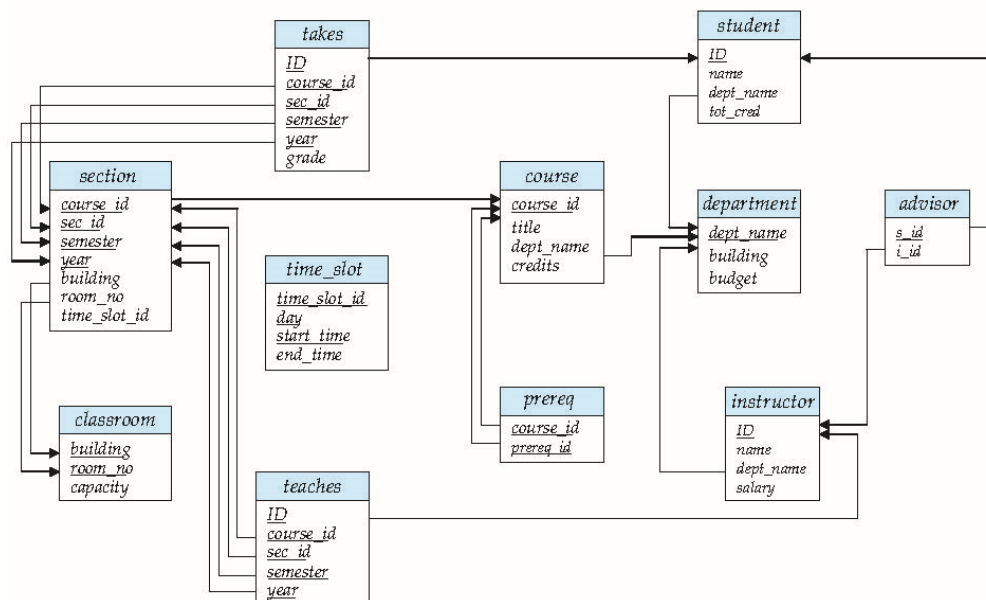
- create table** *student* (  
   *ID*                **varchar**(5),  
   *name*            **varchar**(20) not null,  
   *dept\_name*       **varchar**(20),  
   *tot\_cred*        **numeric**(3,0),  
   **primary key** (*ID*),  
   **foreign key** (*dept\_name*)  
       **references** *department*);
- create table** *takes* (  
   *ID*                **varchar**(5),  
   *course\_id*       **varchar**(8),  
   *sec\_id*           **varchar**(8),  
   *semester*        **varchar**(6),  
   *year*             **numeric**(4,0),  
   *grade*            **varchar**(2),  
   **primary key** (*ID, course\_id, sec\_id, semester, year*) ,  
   **foreign key** (*ID*) **references** *student*,  
   **foreign key** (*course\_id, sec\_id, semester, year*) **references** *section*);
- Note:** *sec\_id* can be dropped from primary key above, to ensure a student cannot be registered for two sections of the same course in the same semester





## And more still

- create table** *course* (  
     *course\_id*     **varchar**(8),  
     *title*        **varchar**(50),  
     *dept\_name*    **varchar**(20),  
     *credits*       **numeric**(2,0),  
     **primary key** (*course\_id*),  
     **foreign key** (*dept\_name*) **references** *department*);





# Updates to tables

- **Insert**
  - **insert into** *instructor* **values** ('10211', 'Smith', 'Biology', 66000);
- **Delete**
  - Remove all tuples from the *student* relation
    - ▶ **delete from** *student*
- **Drop Table**
  - **drop table** *r*
- **Alter**
  - **alter table** *r* **add** *A D*
    - ▶ where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*
    - ▶ All exiting tuples in the relation are assigned *null* as the value for the new attribute
  - **alter table** *r* **drop** *A*
    - ▶ where *A* is the name of an attribute of relation *r*
    - ▶ Dropping of attributes not supported by many databases



PPD

- History of SQL
- Data Definition Language (DDL)
- **Basic Query Structure (DML)**

# BASIC QUERY STRUCTURE





# Basic Query Structure

- A typical SQL query has the form:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

- $A_i$  represents an attribute
  - $R_i$  represents a relation
  - $P$  is a predicate
- The result of an SQL query is a relation



# The select Clause

- The **select** clause lists the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra
- Example: find the names of all instructors:  
**select** *name*  
**from** *instructor*
- NOTE: SQL names are case insensitive (i.e., you may use upper- or lower-case letters.)
  - E.g., *Name*  $\equiv$  *NAME*  $\equiv$  *name*
  - Some people use upper case wherever we use bold font



## The select Clause (Cont.)

- SQL allows duplicates in relations as well as in query results
- To force the elimination of duplicates, insert the keyword **distinct** after select
- Find the department names of all instructors, and remove duplicates

```
select distinct dept_name  
from instructor
```

- The keyword **all** specifies that duplicates should not be removed

```
select all dept_name  
from instructor
```



## The select Clause (Cont.)

- An asterisk in the select clause denotes “all attributes”

**select \***  
**from** *instructor*

- An attribute can be a literal with no **from** clause

**select** '437'

- Results is a table with one column and a single row with value “437”
- Can give the column a name using:

**select** '437' **as** *FOO*

- An attribute can be a literal with **from** clause

**select** 'A'  
**from** *instructor*

- Result is a table with one column and  $N$  rows (number of tuples in the *instructors* table), each row with value “A”



## The select Clause (Cont.)

The **select** clause can contain arithmetic expressions involving the operation, +, −, \*, and /, and operating on constants or attributes of tuples

- The query:

```
select ID, name, salary/12  
from instructor
```

- would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12
- Can rename “*salary/12*” using the **as** clause:

```
select ID, name, salary/12 as monthly_salary
```



# The where Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra
- To find all instructors in Comp. Sci. dept

```
select name  
from instructor  
where dept_name = 'Comp. Sci.'
```

- Comparison results can be combined using the logical connectives **and**, **or**, and **not**
  - To find all instructors in Comp. Sci. dept with salary > 80000

```
select name  
from instructor  
where dept_name = 'Comp. Sci.' and salary > 80000
```

- Comparisons can be applied to results of arithmetic expressions



# The from Clause

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra
- Find the Cartesian product *instructor X teaches*  

```
select *  
from instructor, teaches
```

  - generates every possible instructor – teaches pair, with all attributes from both relations
  - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra)



# Cartesian Product

*instructor*

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

*teaches*

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

Inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Pinance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Pinance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Pinance	90000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...





# Module Summary

- Introduced relational query language
- Familiarized with data definition and basic query structure



# Instructor and TAs

Name	Mail	Mobile
Partha Pratim Das, Instructor	ppd@cse.iitkgp.ernet.in	9830030880
Srijoni Majumdar, TA	majumdarsrijoni@gmail.com	9674474267
Himadri B G S Bhuyan, TA	himadribhuyan@gmail.com	9438911655
Gurunath Reddy M	mgurunathreddy@gmail.com	9434137638

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with “PPD”.



# Database Management Systems

## Module 07: Introduction to SQL/2

**Partha Pratim Das**

*Department of Computer Science and Engineering  
Indian Institute of Technology, Kharagpur*

ppd@cse.iitkgp.ernet.in

**Srijoni Majumdar  
Himadri B G S Bhuyan  
Gurunath Reddy M**



**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
[www.db-book.com](http://www.db-book.com)

Slides marked with 'PPD' are new or edited



# Module Recap

- History of SQL
- Data Definition Language (DDL)
- Basic Query Structure (DML)



# Module Objectives

- To complete the understanding of basic query structure and set operations
- To familiarize with null values and aggregation



# Module Outline

- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions



PPD

- **Additional Basic Operations**
- Set Operations
- Null Values
- Aggregate Functions

# ADDITIONAL BASIC OPERATIONS



# Cartesian Product

- Find the Cartesian product *instructor X teaches*
- - select \***
  - from *instructor, teaches***
    - generates every possible instructor – teaches pair, with all attributes from both relations
    - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra)





# Cartesian Product

*instructor*

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

*teaches*

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

Inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Pinance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Pinance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Pinance	90000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...



# Examples

- Find the names of all instructors who have taught some course and the course\_id

```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID
```

- Equi-Join, Natural Join

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

Inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
<del>10101</del>	<del>Srinivasan</del>	<del>Comp. Sci.</del>	<del>65000</del>	<del>12121</del>	<del>FIN-201</del>	<del>1</del>	<del>Spring</del>	<del>2010</del>
<del>10101</del>	<del>Srinivasan</del>	<del>Comp. Sci.</del>	<del>65000</del>	<del>15151</del>	<del>MU-199</del>	<del>1</del>	<del>Spring</del>	<del>2010</del>
<del>10101</del>	<del>Srinivasan</del>	<del>Comp. Sci.</del>	<del>65000</del>	<del>22222</del>	<del>PHY-101</del>	<del>1</del>	<del>Fall</del>	<del>2009</del>
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
<del>12121</del>	<del>Wu</del>	<del>Finance</del>	<del>90000</del>	<del>10101</del>	<del>CS-101</del>	<del>1</del>	<del>Fall</del>	<del>2009</del>
<del>12121</del>	<del>Wu</del>	<del>Finance</del>	<del>90000</del>	<del>10101</del>	<del>CS-315</del>	<del>1</del>	<del>Spring</del>	<del>2010</del>
<del>12121</del>	<del>Wu</del>	<del>Finance</del>	<del>90000</del>	<del>10101</del>	<del>CS-347</del>	<del>1</del>	<del>Fall</del>	<del>2009</del>
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2010
<del>12121</del>	<del>Wu</del>	<del>Finance</del>	<del>90000</del>	<del>15151</del>	<del>MU-199</del>	<del>1</del>	<del>Spring</del>	<del>2010</del>
<del>12121</del>	<del>Wu</del>	<del>Finance</del>	<del>90000</del>	<del>22222</del>	<del>PHY-101</del>	<del>1</del>	<del>Fall</del>	<del>2009</del>



## Examples

- Find the names of all instructors in the Art department who have taught some course and the course\_id

```
select name, course_id  
from instructor, teaches  
where instructor.ID = teaches.ID and instructor.dept_name = 'Art'
```



# The Rename Operation

- The SQL allows renaming relations and attributes using the **as** clause:

*old-name as new-name*

- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept_name = 'Comp. Sci.'
```

- Keyword **as** is optional and may be omitted  
*instructor as T*  $\equiv$  *instructor T*



## Cartesian Product Example\*

- Relation *emp-super*

<i>person</i>	<i>supervisor</i>
Bob	Alice
Mary	Susan
Alice	David
David	Mary

- Find the supervisor of “Bob”
- Find the supervisor of the supervisor of “Bob”
- Find ALL the supervisors (direct and indirect) of “Bob”



# String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring
  - underscore ( \_ ). The \_ character matches any character
- Find the names of all instructors whose name includes the substring “dar”

```
select name  
from instructor  
where name like '%dar%'
```

- Match the string “100%”

```
like '100 \%' escape '\'
```
- in that above we use backslash (\) as the escape character



## String Operations (Cont.)

- Patterns are case sensitive
- Pattern matching examples:
  - 'Intro%' matches any string beginning with "Intro"
  - '%Comp%' matches any string containing "Comp" as a substring
  - '\_\_\_' matches any string of exactly three characters
  - '\_\_\_ %' matches any string of at least three characters
- SQL supports a variety of string operations such as
  - concatenation (using "||")
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.



# Ordering the Display of Tuples

- List in alphabetic order the names of all instructors  
**select distinct** *name*  
**from** *instructor*  
**order by** *name*
- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
  - Example: **order by** *name desc*
- Can sort on multiple attributes
  - Example: **order by** *dept\_name, name*





## Where Clause Predicates

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq \$90,000$  and  $\leq \$100,000$ )

```
select name  
from instructor  
where salary between 90000 and 100000
```

- Tuple comparison

```
select name, course_id  
from instructor, teaches  
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```



## Duplicates\*

- In relations with duplicates, SQL can define how many copies of tuples appear in the result
- **Multiset** versions of some of the relational algebra operators – given multiset relations  $r_1$  and  $r_2$ :
  1.  $\sigma_{\theta}(r_1)$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$ , and  $t_1$  satisfies selections  $\sigma_{\theta}$ , then there are  $c_1$  copies of  $t_1$  in  $\sigma_{\theta}(r_1)$
  2.  $\Pi_A(r)$ : For each copy of tuple  $t_1$  in  $r_1$ , there is a copy of tuple  $\Pi_A(t_1)$  in  $\Pi_A(r_1)$  where  $\Pi_A(t_1)$  denotes the projection of the single tuple  $t_1$
  3.  $r_1 \times r_2$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$  and  $c_2$  copies of tuple  $t_2$  in  $r_2$ , there are  $c_1 \times c_2$  copies of the tuple  $t_1 \cdot t_2$  in  $r_1 \times r_2$



## Duplicates (Cont.)\*

- Example: Suppose multiset relations  $r_1 (A, B)$  and  $r_2 (C)$  are as follows:

$$r_1 = \{(1, a) (2, a)\} \quad r_2 = \{(2), (3), (3)\}$$

- Then  $\Pi_B(r_1)$  would be  $\{(a), (a)\}$ , while  $\Pi_B(r_1) \times r_2$  would be
- $\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3), (a, 3)\}$
- SQL duplicate semantics:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

is equivalent to the *multiset* version of the expression:

$$\Pi_{A_1, A_2, \dots, A_n} (\sigma_P (r_1 \times r_2 \times \dots \times r_m))$$



- Additional Basic Operations
- **Set Operations**
- Null Values
- Aggregate Functions

# SET OPERATIONS



# Set Operations

- Find courses that ran in Fall 2009 or in Spring 2010  
(**select** *course\_id* **from** *section* **where** *sem* = 'Fall' **and** *year* = 2009)  
**union**  
(**select** *course\_id* **from** *section* **where** *sem* = 'Spring' **and** *year* = 2010)
- Find courses that ran in Fall 2009 and in Spring 2010  
(**select** *course\_id* **from** *section* **where** *sem* = 'Fall' **and** *year* = 2009)  
**intersect**  
(**select** *course\_id* **from** *section* **where** *sem* = 'Spring' **and** *year* = 2010)
- Find courses that ran in Fall 2009 but not in Spring 2010  
(**select** *course\_id* **from** *section* **where** *sem* = 'Fall' **and** *year* = 2009)  
**except**  
(**select** *course\_id* **from** *section* **where** *sem* = 'Spring' **and** *year* = 2010)



## Set Operations (Cont.)

- Find the salaries of all instructors that are less than the largest salary

```
select distinct T.salary  
from instructor as T, instructor as S  
where T.salary < S.salary
```

- Find all the salaries of all instructors

```
select distinct salary  
from instructor
```

- Find the largest salary of all instructors

```
(select "second query" )  
except  
(select "first query")
```



## Set Operations (Cont.)

- Set operations **union**, **intersect**, and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.
- Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:
  - $m + n$  times in  $r$  **union all**  $s$
  - $\min(m, n)$  times in  $r$  **intersect all**  $s$
  - $\max(0, m - n)$  times in  $r$  **except all**  $s$



# NULL VALUES

PPD

- Additional Basic Operations
- Set Operations
- **Null Values**
- Aggregate Functions





# Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist
- The result of any arithmetic expression involving *null* is *null*
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values
  - Example: Find all instructors whose salary is null

```
select name  
from instructor  
where salary is null
```



# Null Values and Three Valued Logic

- Three values – *true*, *false*, *unknown*
- Any comparison with *null* returns *unknown*
  - Example:  $5 < null$  or  $null <> null$  or  $null = null$
- Three-valued logic using the value *unknown*:
  - OR:  $(unknown \text{ or } true) = true$ ,  
 $(unknown \text{ or } false) = unknown$   
 $(unknown \text{ or } unknown) = unknown$
  - AND:  $(true \text{ and } unknown) = unknown$ ,  
 $(false \text{ and } unknown) = false$ ,  
 $(unknown \text{ and } unknown) = unknown$
  - NOT:  $(\text{not } unknown) = unknown$
  - “*P* is unknown” evaluates to true if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*



PPD

- Additional Basic Operations
- Set Operations
- Null Values
- **Aggregate Functions**

# AGGREGATE FUNCTIONS



# Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values



## Aggregate Functions (Cont.)

- Find the average salary of instructors in the Computer Science department  

```
select avg (salary)
from instructor
where dept_name= 'Comp. Sci.';
```
- Find the total number of instructors who teach a course in the Spring 2010 semester  

```
select count (distinct ID)
from teaches
where semester = 'Spring' and year = 2010;
```
- Find the number of tuples in the *course* relation  

```
select count (*)
from course;
```



## Aggregate Functions – Group By

- Find the average salary of instructors in each department

```
select dept_name, avg (salary) as avg_salary
from instructor
group by dept_name;
```

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

dept_name	avg_salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000



## Aggregation (Cont.)

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list

```
/* erroneous query */  
select dept_name, ID, avg (salary)  
from instructor  
group by dept_name;
```



## Aggregate Functions – Having Clause

- Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups





# Null Values and Aggregates

- Total all salaries
  - select** **sum** (*salary*)  
**from** *instructor*
  - Above statement ignores null amounts
  - Result is *null* if there is no non-null amount
- All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
  - **count** returns 0
  - all other aggregates return null



# Module Summary

- Completed the understanding of basic query structure and set operations
- Familiarized with null values and aggregation



PPD

## Instructor and TAs

Name	Mail	Mobile
Partha Pratim Das, Instructor	ppd@cse.iitkgp.ernet.in	9830030880
Srijoni Majumdar, TA	majumdarsrijoni@gmail.com	9674474267
Himadri B G S Bhuyan, TA	himadribhuyan@gmail.com	9438911655
Gurunath Reddy M	mgurunathreddy@gmail.com	9434137638

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with “PPD”.



# Database Management Systems

## Module 08: Introduction to SQL/3

**Partha Pratim Das**

*Department of Computer Science and Engineering  
Indian Institute of Technology, Kharagpur*

ppd@cse.iitkgp.ernet.in

**Srijoni Majumdar  
Himadri B G S Bhuyan  
Gurunath Reddy M**



**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
[www.db-book.com](http://www.db-book.com)

Slides marked with 'PPD' are new or edited



# Module Recap

- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions



# Module Objectives

- To understand nested sub-query in SQL
- To understand processes for data modification



# Module Outline

- Nested Subqueries
- Modification of the Database



PPD

- **Nested Subqueries**
- Modification of the Database

# NESTED SUB-QUERIES





# Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries
- A **subquery** is a **select-from-where** expression that is nested within another query
- The nesting can be done in the following SQL query

```
select  $A_1, A_2, \dots, A_n$   
from  $r_1, r_2, \dots, r_m$   
where  $P$ 
```

as follows:

- $A_i$  can be replaced by a subquery that generates a single value
- $r_i$  can be replaced by any valid subquery
- $P$  can be replaced with an expression of the form:

$B <\text{operation}> (\text{subquery})$

where  $B$  is an attribute and  $<\text{operation}>$  to be defined later



# Subqueries in the Where Clause



## Subqueries in the Where Clause

- A common use of subqueries is to perform tests:
  - For set membership
  - For set comparisons
  - For set cardinality



# Set Membership

- Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
       course_id in (select course_id
                     from section
                     where semester = 'Spring' and year= 2010);
```

- Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
       course_id not in (select course_id
                     from section
                     where semester = 'Spring' and year= 2010);
```



## Set Membership (Cont.)

- Find the total number of (distinct) students who have taken course sections taught by the instructor with *ID* 10101

```
select count (distinct ID)  
from takes  
where (course_id, sec_id, semester, year) in  
      (select course_id, sec_id, semester, year  
       from teaches  
       where teaches.ID= 10101);
```

- Note: Above query can be written in a much simpler manner. The formulation above is simply to illustrate SQL features.



## Set Comparison – “some” Clause

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept name = 'Biology';
```

- Same query using > **some** clause

```
select name  
from instructor  
where salary > some (select salary  
                     from instructor  
                     where dept name = 'Biology');
```



## Definition of “some” Clause\*

- $F <\text{comp}> \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F <\text{comp}> t)$

Where  $<\text{comp}>$  can be:  $<$ ,  $\leq$ ,  $>$ ,  $=$ ,  $\neq$

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$ 
(read: 5 < some tuple in the relation)

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$

However,  $(\neq \text{some}) \not\equiv \text{not in}$



## Set Comparison – “all” Clause

- Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department

```
select name
from instructor
where salary > all (select salary
                     from instructor
                     where dept name = 'Biology');
```





## Definition of “all” Clause\*

- $F \text{ <comp> all } r \Leftrightarrow \forall t \in r (F \text{ <comp> } t)$

$$(5 < \text{all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$$

$$(5 < \text{all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$$

$$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$$

$(\neq \text{all}) \equiv \text{not in}$

However,  $(= \text{all}) \not\equiv \text{in}$



## Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$



## Use of “exists” Clause

- Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
      exists (select *
              from section as T
              where semester = 'Spring' and year = 2010
                  and S.course_id = T.course_id);
```

- **Correlation name** – variable *S* in the outer query
- **Correlated subquery** – the inner query



## Use of “not exists” Clause

- Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name  
from student as S  
where not exists ( (select course_id  
                   from course  
                   where dept_name = 'Biology')  
except  
                  (select T.course_id  
                   from takes as T  
                   where S.ID = T.ID));
```

- First nested query lists all courses offered in Biology
  - Second nested query lists all courses a particular student took
- Note:*  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$
- Note:* Cannot write this query using = **all** and its variants



## Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result
- The **unique** construct evaluates to “true” if a given subquery contains no duplicates
- Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique (select R.course_id
                  from section as R
                  where T.course_id= R.course_id
                     and R.year = 2009);
```



# Subqueries in the From Clause



## Subqueries in the From Clause

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000
  - **select** *dept\_name*, *avg\_salary*  
**from** (**select** *dept\_name*, **avg** (*salary*) **as** *avg\_salary*  
**from** *instructor*  
**group by** *dept\_name*)  
**where** *avg\_salary* > 42000;
- Note that we do not need to use the **having** clause
- Another way to write above query
  - **select** *dept\_name*, *avg\_salary*  
**from** (**select** *dept\_name*, **avg** (*salary*)  
**from** *instructor*  
**group by** *dept\_name*) **as** *dept\_avg* (*dept\_name*, *avg\_salary*)  
**where** *avg\_salary* > 42000;



## With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs
- Find all departments with the maximum budget

```
with max_budget (value) as  
    (select max(budget)  
     from department)  
select department.name  
from department, max_budget  
where department.budget = max_budget.value;
```





## Complex Queries using With Clause\*

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as  
    (select dept_name, sum(salary)  
     from instructor  
     group by dept_name),  
dept_total_avg(value) as  
    (select avg(value)  
     from dept_total)  
select dept_name  
from dept_total, dept_total_avg  
where dept_total.value > dept_total_avg.value;
```



# Subqueries in the Select Clause



## Scalar Subquery

- Scalar subquery is one which is used where a single value is expected
- List all departments along with the number of instructors in each department

```
select dept_name,  
      (select count(*)  
       from instructor  
       where department.dept_name = instructor.dept_name)  
as num_instructors  
from department;
```

- Runtime error if subquery returns more than one result tuple



PPD

- Nested Subqueries
- Modification of the Database**

# MODIFICATION OF THE DATABASE



# Modification of the Database

- Deletion of tuples from a given relation
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation



# Deletion

- Delete all instructors

**delete from** *instructor*

- Delete all instructors from the Finance department

**delete from** *instructor*  
**where** *dept\_name* = 'Finance';

- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building

**delete from** *instructor*  
**where** *dept name* in (**select** *dept name*  
                          **from** *department*  
                          **where** *building* = 'Watson');



## Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor  
where salary < (select avg (salary)  
                  from instructor);
```

- **Problem:** as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
  1. First, compute **avg** (*salary*) and find all tuples to delete
  2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)



# Insertion

- Add a new tuple to *course*

```
insert into course  
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently

```
insert into course (course_id, title, dept_name, credits)  
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to *student* with *tot\_creds* set to null

```
insert into student  
values ('3003', 'Green', 'Finance', null);
```





## Insertion (Cont.)

- Add all instructors to the *student* relation with tot\_creds set to 0

```
insert into student  
select ID, name, dept_name, 0  
from instructor
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation
- Otherwise queries like

```
insert into table1 select * from table1
```

would cause problem



# Updates

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
  - Write two **update** statements:  

```
update instructor  
  set salary = salary * 1.03  
  where salary > 100000;  
update instructor  
  set salary = salary * 1.05  
  where salary <= 100000;
```
  - The order is important
  - Can be done better using the **case** statement (next slide)



## Case Statement for Conditional Updates

- Same query as before but with case statement

```
update instructor  
set salary = case  
    when salary <= 100000 then salary * 1.05  
    else salary * 1.03  
end
```



## Updates with Scalar Subqueries

- Recompute and update `tot_creds` value for all students

**update** *student S*

**set** *tot\_cred* = (**select** **sum**(*credits*)  
                   **from** *takes, course*  
                   **where** *takes.course\_id = course.course\_id and*  
                           *S.ID= takes.ID.and*  
                           *takes.grade <> 'F' and*  
                           *takes.grade is not null*);

- Sets *tot\_creds* to null for students who have not taken any course
- Instead of **sum**(*credits*), use:

**case**  
   **when** **sum**(*credits*) **is not null** **then** **sum**(*credits*)  
   **else** 0  
**end**



# Module Summary

- Introduced nested sub-query in SQL
- Introduced data modification



PPD

## Instructor and TAs

Name	Mail	Mobile
Partha Pratim Das, Instructor	ppd@cse.iitkgp.ernet.in	9830030880
Srijoni Majumdar, TA	majumdarsrijoni@gmail.com	9674474267
Himadri B G S Bhuyan, TA	himadribhuyan@gmail.com	9438911655
Gurunath Reddy M	mgurunathreddy@gmail.com	9434137638

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with “PPD”.



# Database Management Systems

## Module 09: Intermediate SQL/1

**Partha Pratim Das**

*Department of Computer Science and Engineering  
Indian Institute of Technology, Kharagpur*

ppd@cse.iitkgp.ernet.in

**Srijoni Majumdar  
Himadri B G S Bhuyan  
Gurunath Reddy M**



**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
[www.db-book.com](http://www.db-book.com)

Slides marked with 'PPD' are new or edited



# Module Recap

- Nested Subqueries
- Modification of the Database





# Module Objectives

- To learn SQL expressions for Join and Views
- To understand transactions



# Module Outline

- Join Expressions
- Views
- Transactions



PPD

- **Join Expressions**
- Views
- Transactions

# JOIN EXPRESSIONS



# Joined Relations

- **Join operations** take two relations and return as a result another relation
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition).
- It also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the **from** clause



# Types of Join between Relations

- Cross join
- Inner join
  - Equi-join
    - Natural join
- Outer join
  - Left outer join
  - Right outer join
  - Full outer join
- Self-join



# Cross Join

- CROSS JOIN returns the Cartesian product of rows from tables in the join
  - Explicit

```
select *  
from employee cross join department;
```
  - Implicit

```
select *  
from employee, department;
```



## Join operations – Example

- Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that  
prereq information is missing for CS-315 and  
course information is missing for CS-437



# Inner Join



- *course* **inner join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- If specified as **natural**, the 2<sup>nd</sup> *course\_id* field is skipped

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101






# Outer Join

- An extension of the join operation that avoids loss of information
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join
- Uses *null* values



# Left Outer Join



- *course* **natural left outer join** *prereq*


<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101



# Right Outer Join



PPD

- course **natural right outer join** prereq

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101




# Joined Relations

- **Join operations** take two relations and return as a result another relation
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated

<i>Join types</i>	<i>Join Conditions</i>
inner join left outer join right outer join full outer join	natural on <predicate> using ( $A_1, A_1, \dots, A_n$ )



# Full Outer Join



- *course* **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101



## Joined Relations – Examples

- *course* **inner join** *prereq* on  
*course.course\_id = prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- What is the difference between the above (equi\_join), and a natural join?
- *course* **left outer join** *prereq* on  
*course.course\_id = prereq.course\_id*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	<i>null</i>	<i>null</i>



## Joined Relations – Examples

- *course* **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

- *course* **full outer join** *prereq* **using** (*course\_id*)

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101



# VIEWS

PPD

- Join Expressions
- Views**
- Transactions





# Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- A **view** provides a mechanism to hide certain data from the view of certain users
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**



# View Definition

- A view is defined using the **create view** statement which has the form  
  
**create view  $v$  as** < query expression >  
  
▪ where <query expression> is any legal SQL expression
- The view name is represented by  $v$
- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates
- View definition is not the same as creating a new relation by evaluating the query expression
  - Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view



## Example Views

- A view of instructors without their salary  
**create view** *faculty* **as**  
    **select** *ID, name, dept\_name*  
    **from** *instructor*
- Find all instructors in the Biology department  
**select** *name*  
**from** *faculty*  
**where** *dept\_name* = 'Biology'
- Create a view of department salary totals  
**create view** *departments\_total\_salary*(*dept\_name, total\_salary*) **as**  
    **select** *dept\_name, sum (salary)*  
    **from** *instructor*  
    **group by** *dept\_name*;



## Views Defined Using Other Views

- **create view** *physics\_fall\_2009* **as**  
    **select** *course.course\_id, sec\_id, building, room\_number*  
    **from** *course, section*  
    **where** *course.course\_id = section.course\_id*  
          **and** *course.dept\_name = 'Physics'*  
          **and** *section.semester = 'Fall'*  
          **and** *section.year = '2009'*;
- **create view** *physics\_fall\_2009\_watson* **as**  
    **select** *course\_id, room\_number*  
    **from** *physics\_fall\_2009*  
    **where** *building= 'Watson'*;



# View Expansion

- Expand use of a view in a query/another view

```
create view physics_fall_2009_watson as  
(select course_id, room_number  
from (select course.course_id, building, room_number  
        from course, section  
        where course.course_id = section.course_id  
            and course.dept_name = 'Physics'  
            and section.semester = 'Fall'  
            and section.year = '2009')
```

**where** *building* = 'Watson';



## Views Defined Using Other Views

- One view may be used in the expression defining another view
- A view relation  $v_1$  is said to *depend directly* on a view relation  $v_2$  if  $v_2$  is used in the expression defining  $v_1$
- A view relation  $v_1$  is said to *depend on* view relation  $v_2$  if either  $v_1$  depends directly to  $v_2$  or there is a path of dependencies from  $v_1$  to  $v_2$
- A view relation  $v$  is said to be *recursive* if it depends on itself



## View Expansion\*

- A way to define the meaning of views defined in terms of other views
- Let view  $v_1$  be defined by an expression  $e_1$  that may itself contain uses of view relations
- View expansion of an expression repeats the following replacement step:
  - repeat**
    - Find any view relation  $v_i$  in  $e_1$
    - Replace the view relation  $v_i$  by the expression defining  $v_i$
  - until** no more view relations are present in  $e_1$
- As long as the view definitions are not recursive, this loop will terminate



# Recursive View

- In SQL, recursive queries are typically built using these components:
  - A non-recursive seed statement
  - A recursive statement
  - A connection operator
    - The only valid set connection operator in a recursive view definition is UNION ALL
  - A terminal condition to prevent infinite recursion





## Recursive View – Example

- In the context of a relation *flights*:

```
create table flights (
    source    varchar(40),
    destination varchar(40),
    carrier   varchar(40),
    cost      decimal(5,0));
```

source	destination	carrier	cost
Paris	Detroit	KLM	7
Paris	New York	KLM	6
Paris	Boston	American Airlines	8
New York	Chicago	American Airlines	2
Boston	Chicago	American Airlines	6
Detroit	San Jose	American Airlines	4
Chicago	San Jose	American Airlines	2

- Find all the destinations that can be reached from 'Paris'

# Recursive View – Example

```
create recursive view reachable_from (source,destination,depth) as (
  select root.source, root.destination, 0 as depth
  from flights as root
  where root.source = 'Paris'
union all
  select in1.source, out1.destination, in1.depth + 1
  from reachable_from as in1, flights as out1
  where in1.destination = out1.source and
        in1.depth <= 100);
```

- Get the result by simple selection on the view:

```
select distinct source, destination
from reachable_from;
```

source	destination
Paris	Detroit
Paris	New York
Paris	Boston
Paris	Chicago
Paris	San Jose

This example view, *reachable\_from*, is called the *transitive closure* of the *flights* relation

- A non-recursive seed statement
- A recursive statement
- A connection operator
- A terminal condition to prevent infinite recursion

source	destination	carrier	cost
Paris	Detroit	KLM	7
Paris	New York	KLM	6
Paris	Boston	American Airlines	8
New York	Chicago	American Airlines	2
Boston	Chicago	American Airlines	6
Detroit	San Jose	American Airlines	4
Chicago	San Jose	American Airlines	2

Source: [https://info.teradata.com/HTMLPubs/DB\\_TTU\\_16\\_00/index.html#page/SQL\\_Reference%2FB035-1184-160K%2Fsme1472241335807.html%23wwID0EJ23T](https://info.teradata.com/HTMLPubs/DB_TTU_16_00/index.html#page/SQL_Reference%2FB035-1184-160K%2Fsme1472241335807.html%23wwID0EJ23T)



# The Power of Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration
  - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *flights* with itself
    - This can give only a fixed number of levels of reachable destinations
    - Given a fixed non-recursive query, we can construct a database with a greater number of levels of reachable destinations on which the query will not work



# The Power of Recursion

- Computing transitive closure using iteration, adding successive tuples to *reachable\_from*
  - The next slide shows a *flights* relation
  - Each step of the iterative process constructs an extended version of *reachable\_from* from its recursive definition
  - The final result is called the *fixed point* of the recursive view definition.
- Recursive views are required to be **monotonic**. That is, if we add tuples to *flights* the view *reachable\_from* contains all of the tuples it contained before, plus possibly more



## Example of Fixed-Point Computation

source	destination	carrier	cost
Paris	Detroit	KLM	7
Paris	New York	KLM	6
Paris	Boston	American Airlines	8
New York	Chicago	American Airlines	2
Boston	Chicago	American Airlines	6
Detroit	San Jose	American Airlines	4
Chicago	San Jose	American Airlines	2

Iteration #	Tuples in Closure
0	Detroit, New York, Boston
1	Detroit, New York, Boston, San Jose, Chicago
2	Detroit, New York, Boston, San Jose, Chicago



## Update of a View

- Add a new tuple to *faculty* view which we defined earlier

**insert into *faculty* values** ('30765', 'Green', 'Music');

This insertion must be represented by the insertion of the tuple

('30765', 'Green', 'Music', null)

into the *instructor* relation



## Some Updates cannot be Translated Uniquely

- **create view** *instructor\_info* **as**  
    **select** *ID, name, building*  
    **from** *instructor, department*  
    **where** *instructor.dept\_name= department.dept\_name;*
- **insert into** *instructor\_info* **values** ('69987', 'White', 'Taylor');
  - which department, if multiple departments in Taylor?
  - what if no department is in Taylor?
- Most SQL implementations allow updates only on simple views
  - The **from** clause has only one database relation
  - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification
  - Any attribute not listed in the **select** clause can be set to null
  - The query does not have a **group** by or **having** clause



## And Some Not at All

```
create view history_instructors as  
select *  
from instructor  
where dept_name= 'History';
```

- What happens if we insert ('25566', 'Brown', 'Biology', 100000) into *history\_instructors*?





# Materialized Views

- **Materializing a view**: create a physical table containing all the tuples in the result of the query defining the view
- If relations used in the query are updated, the materialized view result becomes out of date
  - Need to **maintain** the view, by updating the view whenever the underlying relations are updated



PPD

- Join Expressions
- Views
- Transactions**

# TRANSACTIONS



# Transactions

- Unit of work
- Atomic transaction
  - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions
- Transactions begin implicitly
  - Ended by **commit work** or **rollback work**
- But default on most databases: each SQL statement commits automatically
  - Can turn off auto commit for a session (e.g. using API)
  - In SQL:1999, can use: **begin atomic .... end**
    - Not supported on most databases



# Module Summary

- Learnt SQL expressions for Join and Views
- Introduced transactions



# Instructor and TAs

Name	Mail	Mobile
Partha Pratim Das, Instructor	ppd@cse.iitkgp.ernet.in	9830030880
Srijoni Majumdar, TA	majumdarsrijoni@gmail.com	9674474267
Himadri B G S Bhuyan, TA	himadribhuyan@gmail.com	9438911655
Gurunath Reddy M	mgurunathreddy@gmail.com	9434137638

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with “PPD”.



# Database Management Systems

## Module 10: Intermediate SQL/2

**Partha Pratim Das**

*Department of Computer Science and Engineering  
Indian Institute of Technology, Kharagpur*

ppd@cse.iitkgp.ernet.in

**Srijoni Majumdar  
Himadri B G S Bhuyan  
Gurunath Reddy M**



**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
[www.db-book.com](http://www.db-book.com)

Slides marked with 'PPD' are new or edited



# Module Recap

- Join Expressions
- Views
- Transactions



# Module Objectives

- To learn SQL expressions for integrity constraints
- To understand more data types in SQL
- To understand authorization in SQL





# Module Outline

- Integrity Constraints
- SQL Data Types and Schemas
- Authorization



PPD

- **Integrity Constraints**
- SQL Data Types and Schemas
- Authorization

# INTEGRITY CONSTRAINTS



# Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency
  - A checking account must have a balance greater than \$10,000.00
  - A salary of a bank employee must be at least \$4.00 an hour
  - A customer must have a (non-null) phone number



## Integrity Constraints on a Single Relation

- **not null**
- **primary key**
- **unique**
- **check (P)**, where P is a predicate



# Not Null and Unique Constraints

- **not null**
  - Declare *name* and *budget* to be **not null**  
*name* **varchar**(20) **not null**  
*budget* **numeric**(12,2) **not null**
- **unique** ( $A_1, A_2, \dots, A_m$ )
  - The unique specification states that the attributes  $A_1, A_2, \dots, A_m$  form a candidate key
  - Candidate keys are permitted to be null (in contrast to primary keys).



# The check clause

- **check** (P)  
where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (  
    course_id varchar (8),  
    sec_id varchar (8),  
    semester varchar (6),  
    year numeric (4,0),  
    building varchar (15),  
    room_number varchar (7),  
    time slot id varchar (4),  
    primary key (course_id, sec_id, semester, year),  
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
);
```



# Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation
  - Example: If “Biology” is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for “Biology”
- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S. A is said to be a **foreign key** of R if for any values of A appearing in R these values also appear in S



## Cascading Actions in Referential Integrity

- **create table** *course* (  
    *course\_id* **char**(5) **primary key**,  
    *title* **varchar**(20),  
    *dept\_name* **varchar**(20) **references** *department*  
)
- **create table** *course* (  
    ...  
    *dept\_name* **varchar**(20),  
    **foreign key** (*dept\_name*) **references** *department*  
        **on delete cascade**  
        **on update cascade**,  
    ...  
)
- alternative actions to cascade: **no action**, **set null**, **set default**





## Integrity Constraint Violation During Transactions

```
create table person (  
  ID char(10),  
  name char(40),  
  mother char(10),  
  father char(10),  
  primary key ID,  
  foreign key father references person,  
  foreign key mother references person)
```

- How to insert a tuple without causing constraint violation ?
  - insert father and mother of a person before inserting person
  - OR, set father and mother to null initially, update after inserting all persons (not possible if father and mother attributes declared to be **not null**)
  - OR defer constraint checking (will discuss later)



PPD

- Integrity Constraints
- **SQL Data Types and Schemas**
- Authorization

# SQL DATA TYPES AND SCHEMAS



## Built-in Data Types in SQL

- **date**: Dates, containing a (4 digit) year, month and date
  - Example: **date** '2005-7-27'
- **time**: Time of day, in hours, minutes and seconds.
  - Example: **time** '09:00:30'      **time** '09:00:30.75'
- **timestamp**: date plus time of day
  - Example: **timestamp** '2005-7-27 09:00:30.75'
- **interval**: period of time
  - Example: interval '1' day
  - Subtracting a date/time/timestamp value from another gives an interval value
  - Interval values can be added to date/time/timestamp values



# Index Creation

```
create table student  
(ID varchar (5),  
 name varchar (20) not null,  
 dept_name varchar (20),  
 tot_cred numeric (3,0) default 0,  
 primary key (ID))
```

- **create index** *studentID\_index* **on** *student*(*ID*)
- Indices are data structures used to speed up access to records with specified values for index attributes

```
select *  
from student  
where ID = '12345'
```

- can be executed by using the index to find the required record, without looking at all records of *student*
- *More on indices in Chapter 11*



# User-Defined Types

- **create type** construct in SQL creates user-defined type

**create type** *Dollars* as numeric (12,2) final

- **create table** *department*  
(*dept\_name* **varchar** (20),  
*building* **varchar** (15),  
*budget* *Dollars*);



# Domains

- **create domain** construct in SQL-92 creates user-defined domain types

```
create domain person_name char(20) not null
```

- Types and domains are similar
- Domains can have constraints, such as **not null**, specified on them

```
create domain degree_level varchar(10)  
constraint degree_level_test  
check (value in ('Bachelors', 'Masters', 'Doctorate'));
```



# Large-Object Types

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*.
  - **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
  - **clob**: character large object -- object is a large collection of character data
  - When a query returns a large object, a pointer is returned rather than the large object itself



- Integrity Constraints
- SQL Data Types and Schemas
- **Authorization**

# AUTHORIZATION





# Authorization

- Forms of authorization on parts of the database:
  - **Read** - allows reading, but not modification of data
  - **Insert** - allows insertion of new data, but not modification of existing data
  - **Update** - allows modification, but not deletion of data
  - **Delete** - allows deletion of data
- Forms of authorization to modify the database schema
  - **Index** - allows creation and deletion of indices
  - **Resources** - allows creation of new relations
  - **Alteration** - allows addition or deletion of attributes in a relation
  - **Drop** - allows deletion of relations



# Authorization Specification in SQL

- The **grant** statement is used to confer authorization  
**grant** <privilege list>  
**on** <relation name or view name> **to** <user list>
- <user list> is:
  - a user-id
  - **public**, which allows all valid users the privilege granted
  - A role (more on this later)
- Granting a privilege on a view does not imply granting any privileges on the underlying relations
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator)



# Privileges in SQL

- **select**: allows read access to relation, or the ability to query using the view
  - Example: grant users  $U_1$ ,  $U_2$ , and  $U_3$  **select** authorization on the *instructor* relation:

**grant select on *instructor* to  $U_1$ ,  $U_2$ ,  $U_3$**

- **insert**: the ability to insert tuples
- **update**: the ability to update using the SQL update statement
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges



# Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization  
**revoke** <privilege list>  
**on** <relation name or view name> **from** <user list>
- Example:  
**revoke select on branch from  $U_1, U_2, U_3$**
- <privilege-list> may be **all** to revoke all privileges the revokee may hold
- If <revokee-list> includes **public**, all users lose the privilege except those granted it explicitly
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation
- All privileges that depend on the privilege being revoked are also revoked



# Roles

- **create role** instructor;  
    **grant instructor to Amit**;
- Privileges can be granted to roles:  
    **grant select on takes to instructor**;
- Roles can be granted to users, as well as to other roles  
    **create role teaching\_assistant**  
    **grant teaching\_assistant to instructor**;
  - *Instructor* inherits all privileges of *teaching\_assistant*
- Chain of roles
  - **create role dean**;
  - **grant instructor to dean**;
  - **grant dean to Satoshi**;



## Authorization on Views

```
create view geo_instructor as  
(select *  
from instructor  
where dept_name = 'Geology');  
grant select on geo_instructor to geo_staff
```

- Suppose that a *geo\_staff* member issues

```
select *  
from geo_instructor;
```
- What if
  - *geo\_staff* does not have permissions on *instructor*?
  - creator of view did not have some permissions on *instructor*?



## Other Authorization Features\*

- **references** privilege to create foreign key
  - **grant reference** (*dept\_name*) **on** *department* **to** Mariano;
  - why is this required?
- transfer of privileges
  - **grant select on** *department* **to** Amit **with grant option**;
  - **revoke select on** *department* **from** Amit, Satoshi **cascade**;
  - **revoke select on** *department* **from** Amit, Satoshi **restrict**;



# Module Summary

- Learnt SQL expressions for integrity constraints
- Familiarized with more data types in SQL
- Discussed authorization in SQL





PPD

## Instructor and TAs

Name	Mail	Mobile
Partha Pratim Das, Instructor	ppd@cse.iitkgp.ernet.in	9830030880
Srijoni Majumdar, TA	majumdarsrijoni@gmail.com	9674474267
Himadri B G S Bhuyan, TA	himadribhuyan@gmail.com	9438911655
Gurunath Reddy M	mgurunathreddy@gmail.com	9434137638

Slides used in this presentation are borrowed from <http://db-book.com/> with kind permission of the authors.

Edited and new slides are marked with “PPD”.